

```
In [1]: import pandas as pd  
import numpy as np
```

```
In [2]: data=pd.read_csv(r"C:\Users\MSI\Downloads\Advertising.csv")  
data
```

```
Out[2]:
```

	TV	Radio	Newspaper	Sales
0	230.1	37.8	69.2	22.1
1	44.5	39.3	45.1	10.4
2	17.2	45.9	69.3	12.0
3	151.5	41.3	58.5	16.5
4	180.8	10.8	58.4	17.9
...
195	38.2	3.7	13.8	7.6
196	94.2	4.9	8.1	14.0
197	177.0	9.3	6.4	14.8
198	283.6	42.0	66.2	25.5
199	232.1	8.6	8.7	18.4

200 rows × 4 columns

```
In [3]: data.head()
```

```
Out[3]:
```

	TV	Radio	Newspaper	Sales
0	230.1	37.8	69.2	22.1
1	44.5	39.3	45.1	10.4
2	17.2	45.9	69.3	12.0
3	151.5	41.3	58.5	16.5
4	180.8	10.8	58.4	17.9

```
In [4]: data.tail()
```

```
Out[4]:
```

	TV	Radio	Newspaper	Sales
195	38.2	3.7	13.8	7.6
196	94.2	4.9	8.1	14.0
197	177.0	9.3	6.4	14.8
198	283.6	42.0	66.2	25.5
199	232.1	8.6	8.7	18.4

```
In [5]: data.shape
```

```
Out[5]: (200, 4)
```


```
In [6]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 200 entries, 0 to 199  
Data columns (total 4 columns):  
#   Column      Non-Null Count  Dtype  
---  -  
0   TV           200 non-null    float64  
1   Radio        200 non-null    float64  
2   Newspaper    200 non-null    float64  
3   Sales        200 non-null    float64  
dtypes: float64(4)  
memory usage: 6.4 KB
```

```
In [7]: import seaborn as sns  
import matplotlib.pyplot as plt
```

```
In [8]: sns.pairplot(data,x_vars=['TV','Radio','Newspaper'],y_vars='Sales',height=7,aspect=0.7,kind='reg')
```

```
Out[8]: <seaborn.axisgrid.PairGrid at 0x17035ff42d0>
```



```
In [9]: feature=['TV','Radio','Newspaper']
```

```
In [10]: x=data[feature]
```

```
In [11]: x=data[['TV','Radio','Newspaper']]
```

```
In [12]: x.head()
```

```
Out[12]:
```

	TV	Radio	Newspaper
0	230.1	37.8	69.2
1	44.5	39.3	45.1
2	17.2	45.9	69.3
3	151.5	41.3	58.5
4	180.8	10.8	58.4

```
In [13]: print(type(x))  
print(x.shape)
```

```
<class 'pandas.core.frame.DataFrame'>  
(200, 3)
```

```
In [14]: y=data['Sales']  
y=data.Sales  
y.head()
```

```
Out[14]: 0    22.1  
1    10.4  
2    12.0  
3    16.5  
4    17.9  
Name: Sales, dtype: float64
```

```
In [15]: print(type(y))  
print(y.shape)
```

```
<class 'pandas.core.series.Series'>  
(200,)
```

```
In [16]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,random_state=1)
```

```
In [17]: print(x_train.shape)
print(x_test.shape)
print(y_train.shape)
print(y_test.shape)
```

```
(150, 3)
(50, 3)
(150,)
(50,)
```

```
In [18]: from sklearn.linear_model import LinearRegression
linreg = LinearRegression()
linreg.fit(x_train,y_train)
```

Out[18]: LinearRegression()

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.
On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [19]: print(linreg.intercept_)
print(linreg.coef_)
```

```
4.633808551125243
[0.05483762 0.10218027 0.00078783]
```

```
In [20]: zip(feature, linreg.coef_)
```

Out[20]: <zip at 0x1702b3fd4c0>

```
In [21]: y_pred = linreg.predict(x_test)
```

```
In [22]: true = [100, 50, 30, 20]
         pred = [90, 50, 50, 30]
```

```
In [23]: print((10 + 0 + 20 + 10) / 4)
         from sklearn import metrics
         print(metrics.mean_absolute_error(true, pred))
```

```
10.0
10.0
```

```
In [24]: import numpy as np
         print((10**2 + 0**2 + 20**2 + 10**2) / 4)
         print(metrics.mean_squared_error(true, pred))
```

```
150.0
150.0
```

```
In [25]: import numpy as np
         print(np.sqrt(((10**2 + 0**2 + 20**2 + 10**2) / 4)))
         print(np.sqrt(metrics.mean_squared_error(true, pred)))
```

```
12.24744871391589
12.24744871391589
```

```
In [26]: print(np.sqrt(metrics.mean_squared_error(y_test, y_pred)))
```

```
1.509610929572584
```

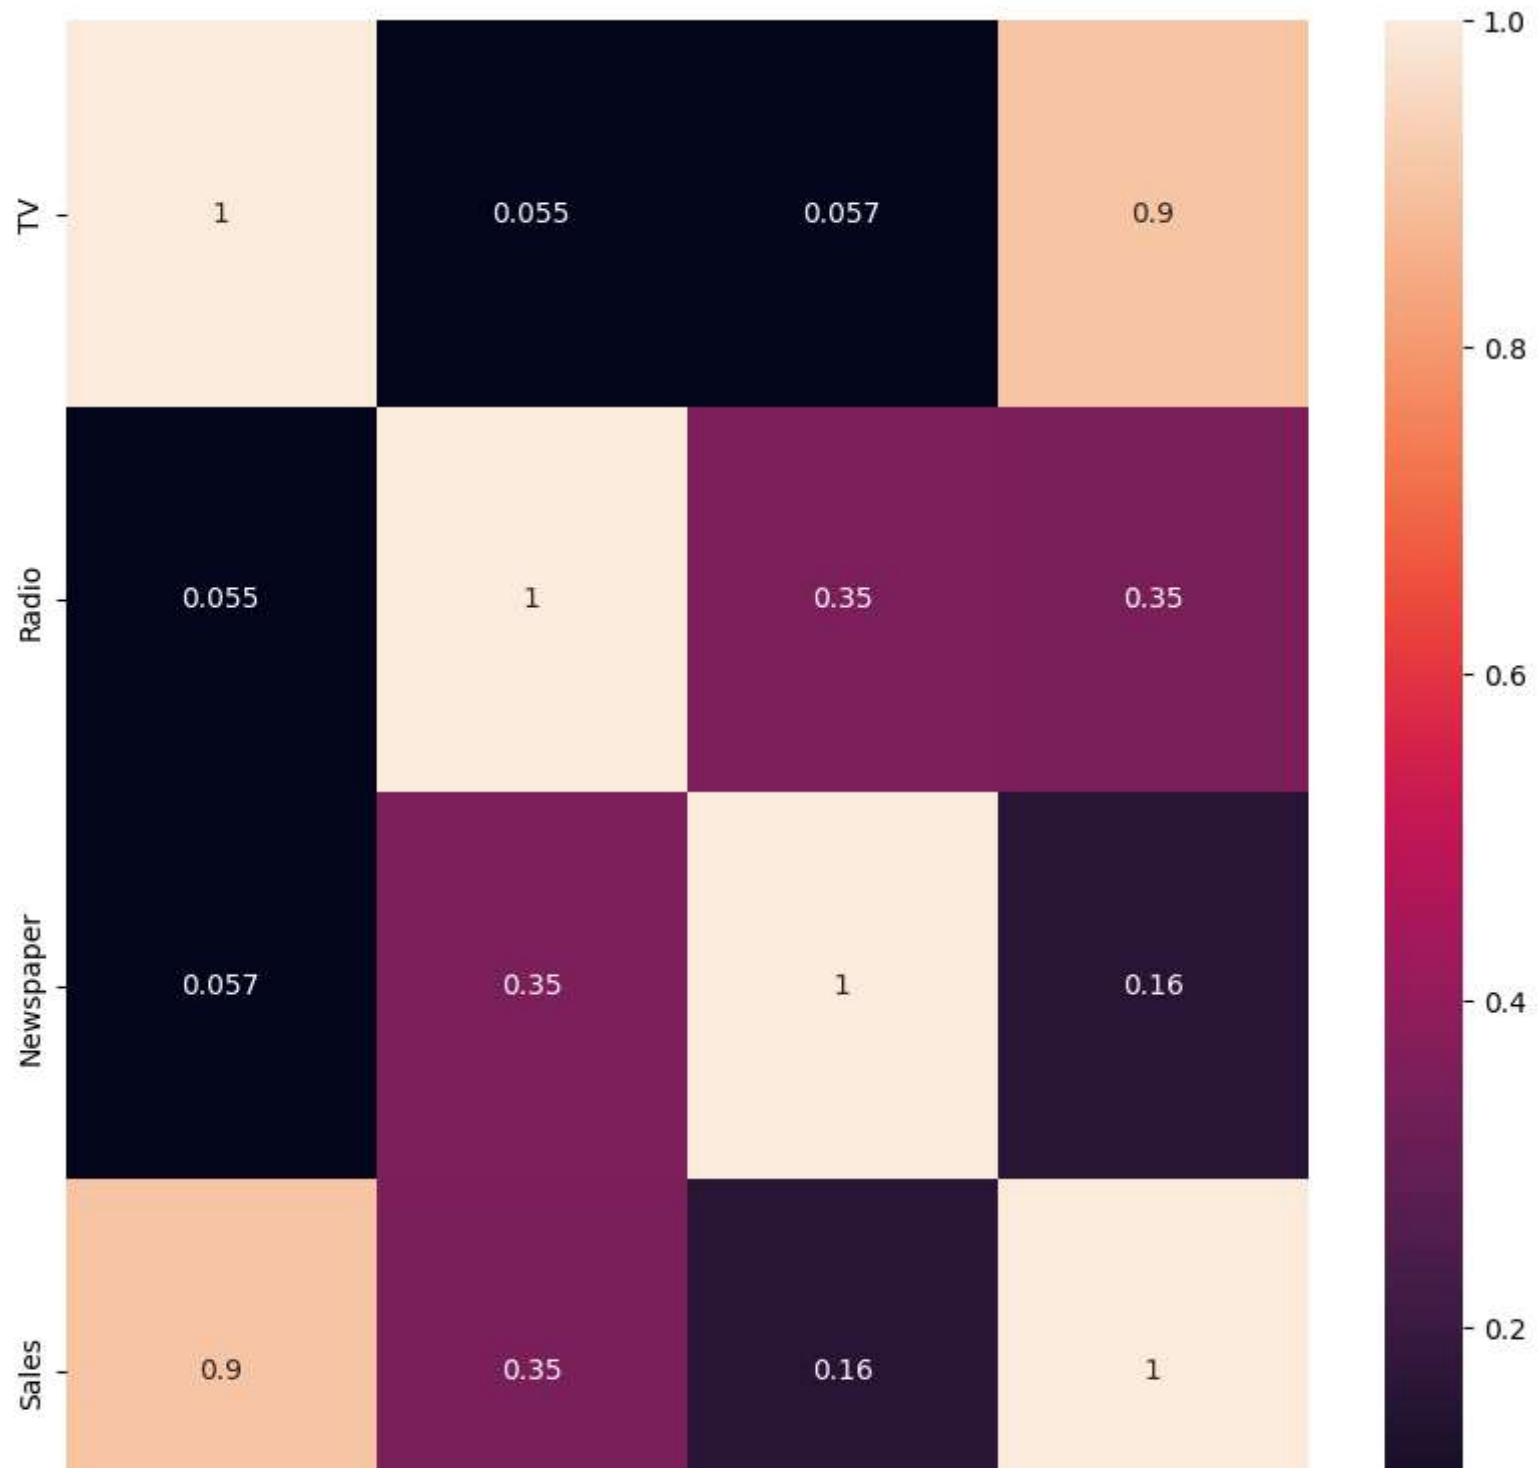
```
In [27]: feature_cols = ['TV', 'Radio']  
x=data[feature_cols]  
y=data.Sales  
x_train,x_test,y_train,y_test = train_test_split(x,y,random_state=1)  
linreg.fit(x_train,y_train)  
y_pred = linreg.predict(x_test)  
print(np.sqrt(metrics.mean_squared_error(y_test,y_pred)))
```

1.5092481618667393

```
In [28]: from sklearn.model_selection import train_test_split  
from sklearn.linear_model import LinearRegression  
from sklearn.linear_model import Ridge, RidgeCV, Lasso  
from sklearn.preprocessing import StandardScaler
```


```
In [29]: plt.figure(figsize = (10, 10))  
sns.heatmap(data.corr(), annot = True)
```

```
Out[29]: <Axes: >
```



```
In [30]: data.drop(columns = ["Radio", "Newspaper"], inplace = True)
sns.pairplot(data)
data.Sales = np.log(data.Sales)
```



```
In [31]: features = data.columns[0:2]
target = data.columns[-1]
x=data[features].values
y=data[target].values
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3,random_state=17)
print("The dimension of x_train is {}".format(x_train.shape))
print("The dimension of x_test is {}".format(x_test.shape))
scaler = StandardScaler()
x_train = scaler.fit_transform(x_train)
x_test = scaler.transform(x_test)
```

The dimension of x_train is (140, 2)

The dimension of x_test is (60, 2)

```
In [32]: lr = LinearRegression()
lr.fit(x_train,y_train)
actual = y_test
train_score_lr = lr.score(x_train,y_train)
test_score_lr = lr.score(x_test,y_test)
print("\nLinear Regression Model:\n")
print("The train score for lr model is {}".format(train_score_lr))
print("The test score for lr model is {}".format(test_score_lr))
```

Linear Regression Model:

The train score for lr model is 1.0

The test score for lr model is 1.0


```
In [33]: ridgeReg = Ridge(alpha=10)
         ridgeReg.fit(x_train,y_train)
         train_score_ridge = ridgeReg.score(x_train, y_train)
         test_score_ridge = ridgeReg.score(x_test, y_test)
         print("\nRidge Model:\n")
         print("The train score for ridge model is {}".format(train_score_ridge))
         print("The test score for ridge model is {}".format(test_score_ridge))
```

Ridge Model:

The train score for ridge model is 0.9902871391941609

The test score for ridge model is 0.984426628514122

```
In [34]: plt.figure(figsize = (10, 10))
plt.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker='*',markersize=5,color='red',label=r'Ridge')
plt.plot(features,lr.coef_,alpha=0.4,linestyle='none',marker='o',markersize=7,color='green',label='LinearReg')
plt.xticks(rotation=90)
plt.legend()
plt.show()
```



```
In [35]: print("\nLasso Model: \n")
lasso = Lasso(alpha = 10)
lasso.fit(x_train,y_train)
train_score_ls =lasso.score(x_train,y_train)
test_score_ls =lasso.score(x_test,y_test)
print("The train score for ls model is {}".format(train_score_ls))
print("The test score for ls model is {}".format(test_score_ls))
```


Lasso Model:

The train score for ls model is 0.0

The test score for ls model is -0.0042092253233847465

```
In [36]: pd.Series(lasso.coef_,features).sort_values(ascending = True).plot(kind = "bar")
```

```
Out[36]: <Axes: >
```

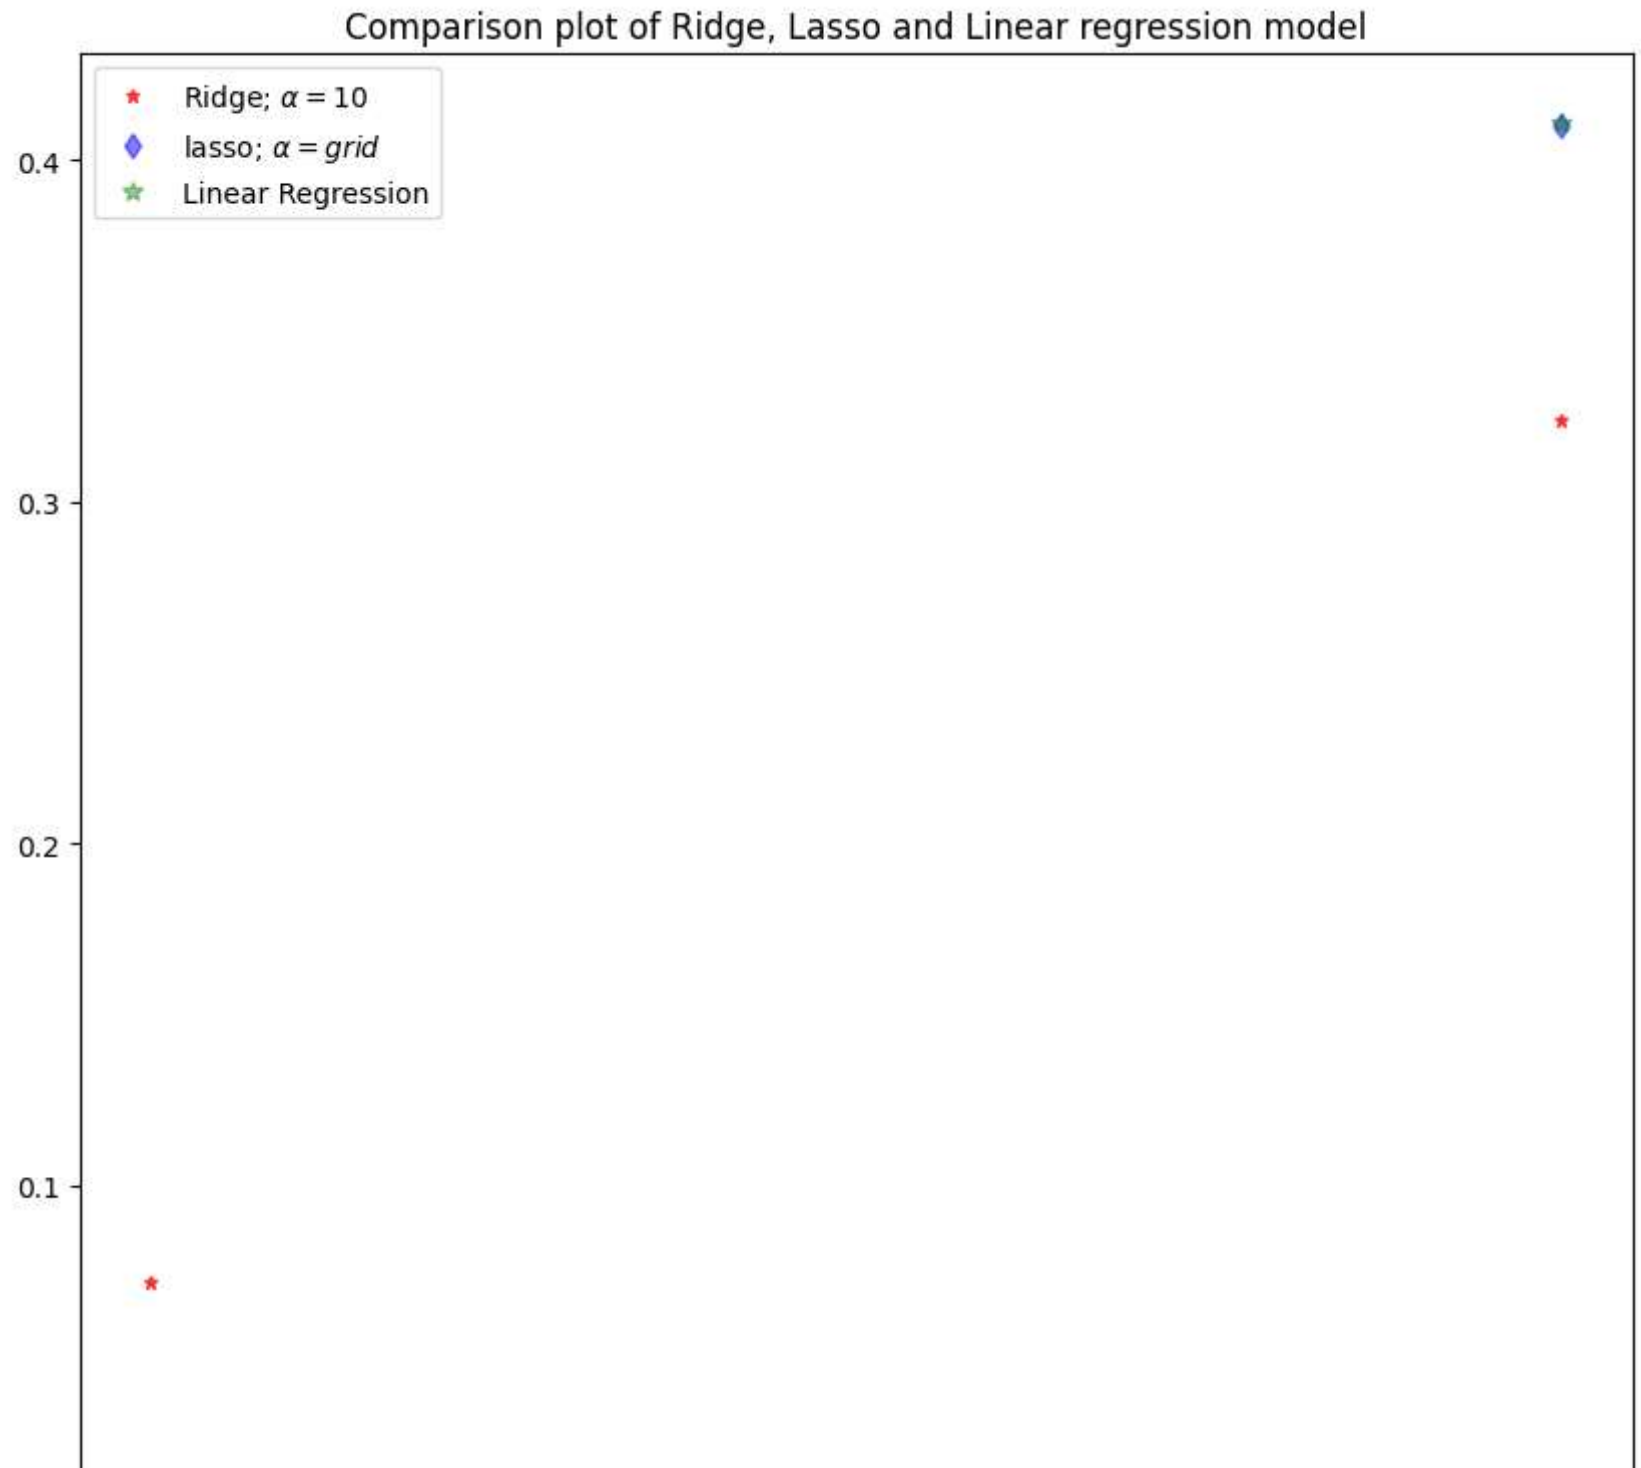


```
In [37]: from sklearn.linear_model import LassoCV
lasso_cv=LassoCV(alphas = [0.0001, 0.001, 0.01, 0.1, 1, 10],random_state=0).fit(x_train,y_train)
print(lasso_cv.score(x_train,y_train))
print(lasso_cv.score(x_test,y_test))
```

```
0.9999999343798134
```

```
0.9999999152638072
```

```
In [38]: plt.figure(figsize=(10, 10))
plt.plot(features,ridgeReg.coef_,alpha=0.7,linestyle='none',marker='*',markersize=5,color='red',label=r'Ridge')
plt.plot(lasso_cv.coef_,alpha=0.5,linestyle='none',marker='d',markersize=6,color='blue',label=r'lasso;  $\lambda$  = 0.5')
plt.plot(features,lr.coef_,alpha=0.4,linestyle='none',marker='*',markersize=7,color='green',label='Linear Regression')
plt.xticks(rotation = 90)
plt.legend()
plt.title("Comparison plot of Ridge, Lasso and Linear regression model")
plt.show()
```



```
In [39]: from sklearn.linear_model import RidgeCV
ridge_cv = RidgeCV(alphas = [0.0001, 0.001, 0.01, 0.1, 1, 10]).fit(x_train, y_train)
print("The train score for ridge model is {}".format(ridge_cv.score(x_train, y_train)))
print("The test score for ridge model is {}".format(ridge_cv.score(x_test, y_test)))
```

The train score for ridge model is 0.99999999997627

The test score for ridge model is 0.999999999962466

```
In [40]: from sklearn.linear_model import ElasticNet
regr = ElasticNet()
regr.fit(x, y)
print(regr.coef_)
print(regr.intercept_)
```

```
[0.00417976 0.          ]
2.026383919311004
```

```
In [41]: y_pred_elastic = regr.predict(x_train)
```

```
In [42]: mean_squared_error = np.mean((y_pred_elastic - y_train)**2)
print("Mean Squared Error on test set", mean_squared_error)
```

Mean Squared Error on test set 0.5538818050142158

```
In [ ]:
```

